

CSC3150 Operating Systems

Project 5 report

Name: Xiao Nan

Student ID: 119010344

Date: 2021/12/07

The Chinese University of Hong Kong, Shen Zhen

0. Introduction:

In homework 5 basic part I implemented a device driver, a kernel module, to control a device which can perform elementary computations. I implemented basic driver operations including device and driver allocation and initialization, reading, writing (with arithmetic routines scheduling) and I/O controlling. In the bonus part I implemented an IRQ handler for the device driver which monitors keyboard interrupts and counts the number of keyboard-generated IRQs happening.

1. Design of the program:

The normal part of the program consists of 5 major functions for driver operations: reading, writing, initializing module, exiting module, and I/O controlling.

In the module initialization function, the program would first register the device by assigning a certain pair of device major and minor number to it using `alloc_chrdev_region` function. Then the device descriptor `cdev` is allocated (using `cdev_alloc()`) and is initialized with the `fops` file operation structure defined in order to map the user-space file operations to the driver. Then after the device is added to the descriptor using `cdev_add()`, the work routine (for temporarily storing job) and the DMA buffer (for storing computational inputs, results, and I/O settings) are allocated (using `kmalloc()`) and initialized (using `memset()`).

In the writing function, the program would first transfer necessary data (operator and operands) to DMA. Then, after initializing the work routine using (`INIT_work()`), the program would schedule works according to the I/O mode read from the DMA buffer: If the I/O mode is blocking then the work would be

put into the working queue using `schedule_work()` and the program would wait till the routine is finished using `flush_scheduled_work()` (note that the readable flag in DMA is set to 1 after the routines are flushed, in order to enable reading the work routine result). If the mode is non-blocking then the program would first set the readable bit in DMA to be 0, indicating that the result in DMA is unreadable, and then return immediately so that CPU can continue doing other works when the work routine is under execution.

The arithmetic function, scheduled in the writing function, is responsible for calculating in a way the DMA suggested using operands and operator. It obtains operator and operands from DMA, judges the operator and performs corresponding calculations in a "switch" sentence. After computation is finished, the result would be stored in the DMA buffer and if the current I/O mode is blocking, the readable flag in DMA would be set 1 so that the result is readable by `ioctl()`.

The `ioctl()` function controls the I/O operation settings. By using such function the user program can set the blocking/non-blocking writing flag in DMA with the value the program desires (using command `HW5_IOCWAITREADABLE`) so as to control writing mode. The function also provides an interface to wait for the readable signal: When receiving command `HW5_SETBLOCK`, the driver would wait in a while loop (by continuously checking the readable flag in DMA) till the flag is set 1 and then return to the user program.

After operations dependent on the driver are completed, the `exit_modules()` would be used to remove the driver and unallocated the device. The `exit_modules()`, registered as

module_exit() function, would first unregister the device and then free the cdev device descriptor. Finally the DMA buffer and work routine would be freed and the module exits successfully.

In the bonus part I added an interrupt handler in the device driver implemented in the normal part. During the module initialization the interrupt handler is registered with IRQ number 1 (corresponding to the keyboard interrupt number) and shared mode (so that it enables multiple handlers for the keyboard interrupt) using request_irq(). In the keyboard interrupt handler function, I simply increment the count stored in DMA (DMACOUNTADDR) with 1 and then return IRQ_HANDLED to notify the OS that the IRQ is handled properly. Finally when exiting the module, the total count in DMA would be printed out. Note that since the interrupt is very sensitive to keyboard hits, the count obtained may be larger than the number of keyboard hits.

2. Runtime environment:

2.1. Linux version:

```
1. namshoo@ubuntu:/$ cat /etc/issue
2. Ubuntu 16.04.5 LTS \n \l
```

2.2. GCC version:

```
1. namshoo@ubuntu:/$ gcc --version
2. gcc (Ubuntu 5.4.0-6ubuntu1~16.04.10) 5.4.0 20160609
3. Copyright (C) 2015 Free Software Foundation, Inc.
4. This is free software; see the source for copying conditions. T
   here is NO
5. warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICU
   LAR PURPOSE.
```

2.3. Linux kernel version:

```
1. namshoo@ubuntu:/$ uname -r
2. 4.15.0
```

3. Steps to execute my program:

3.1. Inserting the driver module:

In command line, you need to first unzip the .zip file. Make sure that the .zip file is under the current directory:

```
1. namshoo@ubuntu:/$ unzip Assignment_5_119010344.zip
```

“cd” to the “Source” execution directory, and type “make” to compile the source code. Make sure that the Makefile, main.c, ioc_hw5.h are under the current directory:

```
1. namshoo@ubuntu:/$ cd Assignment_5_119010344/source/  
2. namshoo@ubuntu:/Assignment_5_119010344/source/$ make
```

3.2. Creating the device file:

Then, check the device MAJOR and MINOR numbers using the command “dmesg | grep chrdev” (MAJOR and MINOR are two numbers identifying a device):

```
1. namshoo@ubuntu:/Assignment_5_119010344/source/$ dmesg | grep chrdev  
2. [xxx.xxxxxx] OS_AS5:init_modules(): register chrdev(MAJOR,MINOR)
```

Then, create the device file using the script “mkdev.sh” with arguments MAJOR and MINOR. Make sure that mkdev.sh is under the current directory:

```
3. namshoo@ubuntu:/Assignment_5_119010344/source/$ sudo ./mkdev.sh MAJOR MINOR
```

3.3. Running the test program:

Then, run the test program, and the outputs would be displayed in terminal:

```
4. namshoo@ubuntu:/Assignment_5_119010344/source/$ ./test
```

3.4 Removing the module and device file:

To remove the device files, you simply need to run the script:

```
1. namshoo@ubuntu:/Assignment_5_119010344/source/$ sudo ./rmdev.sh
```

To remove the driver module and the kernel object files, you simply need to run “make clean” with the help of the Makefile:

```
2. namshoo@ubuntu:/Assignment_5_119010344/source/$ make clean
```

Then, the kernel message outputs would be printed out in the terminal screen (due to the implementation in the Makefile).

4. Program output screenshots:

The results for test program (the case with all arithmetic functions tested) in terminal and in kernel dmesg are shown in figure 4.1, 4.2, respectively. The kernel dmesg output corresponding to the demo case is shown in figure 4.3:

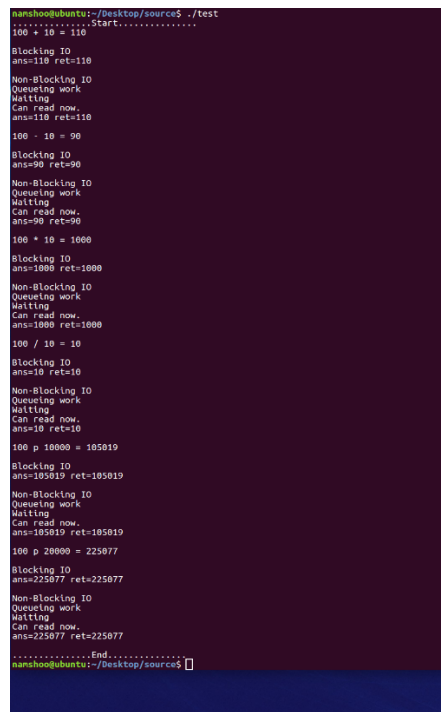
A terminal window with a dark purple background and white text. The prompt is 'namshoo@ubuntu:~/desktop/source\$./test'. The output shows a series of arithmetic operations: addition (100 + 10 = 110), subtraction (100 - 10 = 90), multiplication (100 * 10 = 1000), and division (100 / 10 = 10). Each operation is followed by a 'Blocking IO' section with 'ans' and 'ret' values, and a 'Non-Blocking IO' section with 'Queueing work', 'Waiting', 'Can read now.', and 'ans'/'ret' values. The final output shows a large multiplication result: '100 p 10000 = 105019' and '100 p 20000 = 225077'. The terminal ends with 'End.....' and the prompt 'namshoo@ubuntu:~/desktop/source\$'.

Figure 4.1: Terminal results

```

5705.253089] OS_ASS: init_modules(): .....Start.....
5705.253092] OS_ASS: init_modules(): register chrdev(243,1)
5705.253113] OS_ASS: init_modules(): request_irq 1 returns 0
5705.253113] OS_ASS: init_modules(): allocate dma buffer
5725.585245] OS_ASS: drv_open(): device open
5725.585250] OS_ASS: drv_ioctl(): My STUID is = 123456789
5725.585252] OS_ASS: drv_ioctl(): RW OK
5725.585253] OS_ASS: drv_ioctl(): IOC OK
5725.585255] OS_ASS: drv_ioctl(): IRQ OK
5725.585297] OS_ASS: drv_ioctl(): Blocking IO
5725.585300] OS_ASS: drv_write(): queue work
5725.585304] OS_ASS: drv_write(): block
5725.589596] OS_ASS: drv_arithmetic_routine(): 100 + 10 = 110
5725.589633] OS_ASS: drv_read(): ans = 110
5725.589695] OS_ASS: drv_ioctl(): Non-Blocking IO
5725.589701] OS_ASS: drv_write(): queue work
5725.589729] OS_ASS: drv_arithmetic_routine(): 100 + 10 = 110
5725.589741] OS_ASS: drv_ioctl(): wait readable 1
5725.589745] OS_ASS: drv_read(): ans = 110
5725.589754] OS_ASS: drv_ioctl(): Blocking IO
5725.589756] OS_ASS: drv_write(): queue work
5725.589758] OS_ASS: drv_write(): block
5725.589762] OS_ASS: drv_arithmetic_routine(): 100 - 10 = 90
5725.589794] OS_ASS: drv_read(): ans = 90
5725.589833] OS_ASS: drv_ioctl(): Non-Blocking IO
5725.589838] OS_ASS: drv_write(): queue work
5725.589842] OS_ASS: drv_arithmetic_routine(): 100 - 10 = 90
5725.589848] OS_ASS: drv_ioctl(): wait readable 1
5725.589874] OS_ASS: drv_read(): ans = 90
5725.589883] OS_ASS: drv_ioctl(): Blocking IO
5725.589885] OS_ASS: drv_write(): queue work
5725.589887] OS_ASS: drv_write(): block
5725.589890] OS_ASS: drv_arithmetic_routine(): 100 * 10 = 1000
5725.589907] OS_ASS: drv_read(): ans = 1000
5725.589938] OS_ASS: drv_ioctl(): Non-Blocking IO
5725.589941] OS_ASS: drv_write(): queue work
5725.589946] OS_ASS: drv_arithmetic_routine(): 100 * 10 = 1000
5725.589972] OS_ASS: drv_ioctl(): wait readable 1
5725.589977] OS_ASS: drv_read(): ans = 1000
5725.590007] OS_ASS: drv_ioctl(): Blocking IO
5725.590009] OS_ASS: drv_write(): queue work
5725.590011] OS_ASS: drv_write(): block
5725.590014] OS_ASS: drv_arithmetic_routine(): 100 / 10 = 10
5725.590029] OS_ASS: drv_read(): ans = 10
5725.590058] OS_ASS: drv_ioctl(): Non-Blocking IO
5725.590062] OS_ASS: drv_write(): queue work
5725.590066] OS_ASS: drv_arithmetic_routine(): 100 / 10 = 10
5725.590093] OS_ASS: drv_ioctl(): wait readable 1
5725.590098] OS_ASS: drv_read(): ans = 10
5726.392087] OS_ASS: drv_ioctl(): Blocking IO
5726.392090] OS_ASS: drv_write(): queue work
5726.392093] OS_ASS: drv_write(): block
5726.995059] OS_ASS: drv_arithmetic_routine(): 100 p 10000 = 105019
5726.995136] OS_ASS: drv_read(): ans = 105019
5726.995153] OS_ASS: drv_ioctl(): Non-Blocking IO
5726.995156] OS_ASS: drv_write(): queue work
5726.995159] OS_ASS: drv_ioctl(): wait readable 1
5727.627251] OS_ASS: drv_arithmetic_routine(): 100 p 10000 = 105019
5727.635886] OS_ASS: drv_read(): ans = 105019
5731.072716] OS_ASS: drv_ioctl(): Blocking IO
5731.072720] OS_ASS: drv_write(): queue work
5731.072723] OS_ASS: drv_write(): block
5733.674038] OS_ASS: drv_arithmetic_routine(): 100 p 20000 = 225077
5733.674136] OS_ASS: drv_read(): ans = 225077
5733.674155] OS_ASS: drv_ioctl(): Non-Blocking IO
5733.674158] OS_ASS: drv_write(): queue work
5733.674914] OS_ASS: drv_ioctl(): wait readable 1
5736.253867] OS_ASS: drv_arithmetic_routine(): 100 p 20000 = 225077
5736.266833] OS_ASS: drv_read(): ans = 225077
5736.266954] OS_ASS: drv_release(): device close
5758.486527] OS_ASS: exit_modules(): interrupt count = 209
5758.486542] OS_ASS: exit_modules(): free dma buffer
5758.486542] OS_ASS: exit_modules(): unregister chrdev
5758.486575] OS_ASS: exit_modules(): .....End.....
hamshoo@ubuntu: ~/Desktop/source$

```

Figure 4.2: dmesg results (full)

```

2007.524329] OS_ASS: init_modules(): .....Start.....
2007.524331] OS_ASS: init_modules(): register chrdev(243,1)
2007.524336] OS_ASS: init_modules(): request_irq 1 returns 0
2007.524336] OS_ASS: init_modules(): allocate dma buffer
2020.075013] OS_ASS: drv_open(): device open
2020.075020] OS_ASS: drv_ioctl(): My STUID is = 123456789
2020.075022] OS_ASS: drv_ioctl(): RW OK
2020.075024] OS_ASS: drv_ioctl(): IOC OK
2020.075025] OS_ASS: drv_ioctl(): IRQ OK
2023.471090] OS_ASS: drv_ioctl(): Blocking IO
2023.471095] OS_ASS: drv_write(): queue work
2023.471098] OS_ASS: drv_write(): block
2026.070887] OS_ASS: drv_arithmetic_routine(): 100 p 20000 = 225077
2026.070995] OS_ASS: drv_read(): ans = 225077
2026.071012] OS_ASS: drv_ioctl(): Non-Blocking IO
2026.071015] OS_ASS: drv_write(): queue work
2026.071017] OS_ASS: drv_ioctl(): wait readable 1
2028.725280] OS_ASS: drv_arithmetic_routine(): 100 p 20000 = 225077
2028.730283] OS_ASS: drv_read(): ans = 225077
2028.730363] OS_ASS: drv_release(): device close
2044.102066] OS_ASS: exit_modules(): interrupt count = 134
2044.102105] OS_ASS: exit_modules(): free dma buffer
2044.102105] OS_ASS: exit_modules(): unregister chrdev
2044.102113] OS_ASS: exit_modules(): .....End.....
hamshoo@ubuntu: ~/Desktop/source$

```

Figure 4.3: dmesg results (part, corresponding to demo)

5. Things I've learnt from the project:

In this project I learnt about how Linux manages the devices by mapping device into kernel space data structures and performing operations based on kernel functions. I learnt how Linux can read and write data with the help of the device and how DMA can accelerate data reading and writing and help improve the CPU utilization rate. I also learnt that Linux uses the mechanism of working queue to schedule the works for devices and for synchronizing blocking and non-blocking behaviors. I learnt what a device driver is and how that can be implemented in Linux kernel (especially how the logical structure of the device and its corresponding setting can be represented in the kernel). Moreover, from the bonus part, I learnt how the interrupt handler works when an interrupt is taking place. I learnt the way to register an interrupt handler for an IRQ and make it work in kernel.